

Rules for fixed checks. AI for meaning.

Make the choice at each step, not once for the whole inbox.




USE RULES

Can a verified field decide the action?

 TRUSTED FORM

Request type: Billing

→ **Use the owner lookup**



USE AN LLM

Does context change the answer?

 EMAIL EXAMPLE

“No rush. Launch is tomorrow.
Access is blocked.”

→ **Extract impact + deadline**



KEEP HUMAN REVIEW

Is context missing or approval required?

 EMAIL EXAMPLE

“Any update on my earlier request?”

→ **Review before acting**

Decision guide. Email examples are illustrative, not interview quotes.

AI reads the email. Rules run the workflow.

A controlled handoff, with one model call where interpretation is needed.




Failed checks or missing context → manual review. No auto-send.

Proposed enterprise pilot, adapted from the email-assistant workflow.

Does the LLM actually help?

Compare both approaches on the same labelled emails.



 Right owner

 Missed urgency

 BA minutes

 Cost + latency

 Keep AI only when quality holds and total BA effort falls.

Test tools: Gmail sample • Sheets scorecard • Make run history.

What does the LLM add?


Illustrative model-token cost for 1,000 emails. USD.



Rules only

\$0

LLM token charge
No model calls



Rules + GPT-4o mini

\$0.41

LLM token charge, approximately
1,000 model calls

Per call: 1,500 input + 300 output tokens. One call per email; no retries.

TOTAL OPERATING COST



Model tokens

+



Make usage

+



Support + review

Sources: OpenAI GPT-4o mini pricing and Make Credits. Checked 05 Sep 2026.

The tools exist. Live access needs approval.

Gmail and Make were in place. A corporate LLM plan was possible.

IN PLACE



Gmail inbox
Team mail already in use



Make.com
Automation platform available

PREPARE



Email sample
Clean text + labelled cases



Routing policy
Categories + approved owners

APPROVE



Corporate API
Plan possible; setup pending



Data access
Permissions + retention

 **Approved samples first. Live inbox after sign-off.**

Corporate chat access and API billing are separate. Source: OpenAI billing; readiness: project context.

Email triage is a practical first pilot.

Target the launch-related backlog before work reaches technical teams.




MY ASSESSMENT

High

technical feasibility

Existing stack.
Small scope.
Manual fallback.



Recommendation: a supervised pilot. Feasibility is an assessment, not a completed deployment result.