

Support-email triage

Reduce repetitive intake while keeping the assistant in control.

PROPOSED RELEASE

One inbox.
A clear next action.
A person decides.

PRIMARY USER

The assistant / business analyst responsible for support-email intake and review.

One approved Gmail queue. Suggestions and drafts, not automatic replies or autonomous resolution.

REVIEW CARD

Synthetic example

Service unavailable before tomorrow's event

Service issue

Urgent

Suggested owner

IT support

Reason

Tomorrow's event is blocked.

 Pending assistant review

MEASURE Handling effort • routing quality • urgent misses

Basis: Project Context & Mandate, pp. 1–2; Business Discovery (Final), pp. 3 and 9. Example labels are illustrative.

Core requirements

The first release turns an email into a traceable suggestion, not an autonomous action.



01

Controlled intake

Use one approved inbox. Capture the message ID; prevent duplicate processing.



02

Classify and prioritize

Suggest type, priority, owner and reason. Apply agreed rules before AI judgment.



03

Draft safely

Use approved acknowledgement wording. Never invent a fix, date or commitment.



04

Keep review in control

Surface urgent and uncertain cases. The assistant can accept, edit or escalate.



05

Keep an audit trail

Log the source, output, prompt/model version, reviewer changes and run status.



06

Recover without losing work

Validate outputs; flag failures. Keep the original inbox available for manual handling.


NOT IN SCOPE

Auto-send · autonomous resolution · unrestricted attachments · enterprise-wide rollout

Basis: Business Discovery (Final), p. 9; reference automation assignment, pp. 1–2. Requirements are proposed for this enterprise workflow.

Proposed Make.com workflow

Automate preparation and visibility. Keep the assistant’s review before the final action.



Assistant reviews Accept / edit / escalate → approved handoff or manual reply → update log

Invalid outputs and failed runs are flagged for review. No automatic external replies; the inbox remains the fallback.

Basis: Business Discovery (Final), p. 9; reference prototype, p. 1. Queue names and the scenario layout are illustrative, not a live screenshot.

User stories and acceptance criteria

Primary user: the assistant who reviews the inbox. These checks define the proposed release.

ID	AS THE ASSISTANT, I WANT...	ACCEPTED WHEN...
01	clear triage suggestions, so I can review faster.	Required fields and a reason are present. Labels use the agreed category and owner list; suggestions match agreed test expectations.
02	urgent issues highlighted, so I do not miss critical work.	Urgent challenge cases are flagged. Mixed-topic or missing-history emails go to review rather than a guessed route.
03	a safe, editable draft, so replies take less effort.	The draft uses approved wording without unsupported promises. No external reply is sent automatically.
04	control of the next action, so I can correct suggestions.	I can accept, edit or escalate. The reviewed decision and source ID are logged before the final handoff.
05	reliable, traceable processing, so no email is lost or repeated.	Replaying a message creates no second handoff. Invalid output or a failed run is flagged, and the original inbox item remains available.

Before pilot: agree routing-quality and review-effort targets. A small test set does not establish production accuracy.

Proposed acceptance checks derived from Business Discovery (Final), pp. 9–10, and the reference prototype’s documented limitations.

Start with a golden dataset

A labelled set of emails and expected answers, checked by the assistant before testing.



ID	SAMPLE EMAIL / TEST INPUT	EXPECTED RESULT	REVIEW CHECK
E01	"Please give me dashboard access next week."	Access · routine IT support suggestion	Routine review
E02	"No rush, but the service is down and tomorrow's event cannot run."	Service issue · urgent IT support suggestion	Urgency is not based on tone
E03	"Sensitive records were sent to the wrong recipient."	Privacy risk · urgent Security review	Restrict visibility; do not broadcast
E04	"I need dashboard access and have a question about an invoice."	Multiple topics Owner needs review	Assistant confirms all relevant owners
E05	"Any update on my earlier request?"	Missing context No guessed owner	Request context or review history

Synthetic labels require user approval. Keep prompt-tuning examples separate; also test duplicates, invalid outputs and failed runs.

The raw 25-ticket reference dataset was not attached. Basis for the method: Business Discovery (Final), p. 10; reference prototype, pp. 1–2.

Delivery plan: validate before live integration

Dataset first, demo second, user sign-off third. Connect the final Make.com workflow only after approval.

01



Golden dataset

ASSISTANT + BUILDER

Agree labels and owners.
Create sample emails.
Keep tuning and test examples separate.

EXIT CHECK

Labels and test set
approved by assistant.

02



Demo prototype

BUILDER

Create demo accounts:
Gmail, Sheets and Slack.
Test rules and OpenAI
on the sample emails.

EXIT CHECK

Frozen-set results
and failures recorded.

03



User validation

ASSISTANT + SERVICE OWNERS

Review routes and drafts.
Test urgent cases.
Check review effort.
Fix issues and retest.

EXIT CHECK

Users approve behavior
and team ownership.

04



Final Make.com integration

BUILDER + AUTHORIZED IT

Build final Make scenario.
Use approved accounts.
Run without live actions.
Then pilot one queue.

EXIT CHECK

Access, fallback and
runbook checks passed.

Live-use gate: user validation + IT approval + tested fallback. Then monitor effort, quality, failed runs and cost.