

Defining AI-assisted email triage

FROM 03 · OPPORTUNITY ASSESSMENT

Selected workflow: support-email triage.

In the enterprise IT team, the business analyst handled support emails manually, from reading and classification to routing and replies. I defined an AI-assisted workflow to reduce that preparation work without removing human review.

WORK I SET OUT TO REDUCE

The analyst read each message, interpreted the request, set urgency, found the owner and wrote a reply.

SOLUTION I DEFINED

The workflow prepares the type, priority, owner suggestion and draft. The assistant reviews the result and sends the reply.

MY STARTING POINT

I used my working onboarding-email prototype as the reference pattern for this enterprise solution.

Sources: Project Context & Mandate, pp. 1–2; Business Discovery (Final), pp. 3, 8–9; reference automation summary, p. 1.

What I included in the first version

I limited the scope to one inbox, structured triage, team routing and draft replies.

REQUIREMENT	BEHAVIOR DEFINED FOR THE FIRST VERSION
01 Email intake	A new email starts the workflow from one Gmail inbox; the message content is passed to the AI step.
02 Triage and routing	Each result includes a request type, priority, route and reason. Slack channels make the handoff visible.
03 Draft replies	The workflow generates a draft for human review, without inventing a resolution or commitment.
04 Human control	The assistant reviews replies and ambiguous cases. No customer response is sent automatically.
05 Decision logging	The structured JSON result is parsed and saved in Google Sheets before it is routed to Slack.

NOT INCLUDED

Auto-send, autonomous resolution or CRM conversation-history integration.

Sources: Reference automation summary, pp. 1–2; Business Discovery (Final), p. 9.

User stories and acceptance criteria

I used the human assistant’s tasks to define what the first version had to demonstrate.

ID	USER STORY	ACCEPTANCE CRITERIA
01	As the assistant, I want type and owner suggestions so I can route requests faster.	The output contains a valid type, priority, route and reason. A labelled test email reaches its expected Slack channel.
02	As the assistant, I want urgent requests highlighted so I can act before a deadline is missed.	A request with an imminent deadline stays urgent even when politely worded; urgent cases reach the escalation channel.
03	As the assistant, I want a draft reply so I spend less time writing from scratch.	A draft is available for review, contains no unsupported promise and is never sent automatically.
04	As the assistant, I want unclear or mixed requests flagged so important context is not lost.	Missing-history emails go to human review. Multi-topic messages retain the relevant issues and reach the appropriate teams.
05	As the assistant, I want a decision log so I can trace and correct a misrouted request.	The parsed result is recorded in Sheets before routing, with classification, reason and confidence available for review.

Sources: Reference automation summary, pp. 1–2; Business Discovery (Final), pp. 9–10.

How I prepared the golden dataset

After mapping the specialist’s workflow, I reviewed 25 onboarding tickets to define request types and expected outputs for testing.

ILLUSTRATIVE ENTRIES · Sample email + expected outcome + behavior to check

ID	SAMPLE EMAIL	EXPECTED OUTCOME	TEST FOCUS
E01	"Can you help us complete account setup?"	Onboarding	Basic classification
E02	"No rush, but we launch in 48 hours and setup is blocked."	Urgent escalation Priority P1	Urgency, not tone
E03	"We need setup help and have an invoice question."	Onboarding + sales/billing	More than one issue
E04	"Any update on my earlier request?"	Specialist review	Missing history

PROTOTYPE TESTING

I tested 15 tickets across request types and priorities.

Source: Reference automation summary, pp. 1–2. Sample emails are illustrative, not original ticket text.

My delivery sequence

I structured delivery in four stages, keeping user validation before final Make.com integration.

DOCUMENTED REFERENCE PROTOTYPE		NEXT STAGES IN THE DELIVERY PLAN	
01 Golden dataset I reviewed the tickets, grouped request types and defined expected outputs. Evidence: taxonomy + test set	02 Demo prototype I used demo accounts to test classification, routing and drafts on sample tickets. Evidence: demo + test results	03 User validation The assistant reviews routes and drafts, flags errors and confirms the revised behavior. Gate: user approval + retest	04 Final Make.com integration The validated workflow connects to approved accounts, then starts with one live queue. Gate: access approval + fallback checks

Final integration follows user approval, with approved access and a tested manual fallback.

Sources: Reference automation summary, pp. 1–2; Business Discovery (Final), pp. 9–10.