



One AI step. Fixed handoffs. Human-owned replies.

A bounded email assistant: interpret the request, record the decision and notify the right people.



Build: assignment summary, pp. 1-2. Enterprise reviewer role adapted from the project context.

The actual Make.com workflow



Original screenshot supplied by the author. | [Routing check: Make router documentation](#)

 **01 Interpret once**
One OpenAI step returns classification, priority and a draft.

 **02 Log, then notify**
The parsed decision is written to Sheets before Slack routing.

 **03 Stop before sending**
Slack messages stay internal. No customer-send module is shown.

 **Check fallback routing**
Test general fallback with draft review enabled.

I kept the stack small enough to inspect and change.

The choices balance reasoning quality, cost, integration effort, control and maintainability.



Make.com over custom orchestration

Native connectors and run logs made the first version fast to build and inspect.

Trade-off: platform dependence, usage credits and less control over complex recovery.



A compact OpenAI model

One call handles classification, priority and a draft. Cost and reasoning drove the choice.

Trade-off: test urgency and draft quality alongside cost and latency. Price alone is not enough.



Rules around the LLM

The model interprets meaning. Make parses, logs and routes to configured destinations.

Trade-off: predictable paths, but no autonomous tool choice or missing-history resolution.



Sheets + Slack for the prototype

A visible decision log and familiar review channel made the prototype easy to inspect.


Trade-off: editable logs and chat lack durable audit records and ticket-lifecycle controls.

Not in v1: CRM memory, knowledge retrieval (RAG), fine-tuning or multi-agent planning.

Confirm the exact API model ID in Make before release. | [Decision lens: OpenAI model selection](#)

Control the data path, not just the API key.

Pilot access design: the tools are available; permission to process live email is a separate gate.



OpenAI returns its output to Make. Only Make writes to Sheets and Slack; the model has no direct send action.

Least-privilege access
Approve the inbox and inspect actual Gmail scopes. Limit Make editors, API keys and Slack channel membership.

Minimum necessary data
Send only needed text. Exclude unnecessary sensitive fields and attachments. Avoid copying full emails into shared logs.

Corporate API + retention
Use a company-owned API project with separate billing. Approve retention across OpenAI, Make, Sheets and Slack.

[Gmail scopes](#) | [OpenAI API billing](#) | [OpenAI data controls](#) | Requirements for live use, not verified connector settings.

Automate internal handoffs. Keep replies human.

The specialist reviews, edits and sends in Gmail, outside this Make scenario.

ILLUSTRATIVE REVIEW VIEW

 **#specialist_draft-reply**

INBOUND EMAIL


**School launches in 48 hours.
Setup is still blocked.**

P1 / LAUNCH BLOCKER

Suggested draft

Thanks for flagging this.
Which setup step is blocked?

Review the source email, facts, recipient, tone and any commitments before sending.



INTERNAL ROUTING IS AUTOMATED

 **Onboarding / Sales & billing / Support & product**

 **P0/P1, privacy or safety → escalation**
Unknown request or missing history → specialist review

Current version: manual review. Slack approval buttons and automated customer sending are not implemented.

Build: assignment summary, pp. 1-2. | [Review basis: OpenAI safety guidance](#)

Guardrails are more than instructions in a prompt.

Behavior rules, output checks, restricted actions and traceability each address a different failure mode.

IMPLEMENTED IN THE PROTOTYPE



Behavior guardrail

Defined request scope. Urgency follows deadlines and impact, not the sender's tone.



Output-format guardrail

JSON is parsed before results are logged and routed.



Routing + human-control guardrail

Urgent and general-review paths. Drafts stay internal until the specialist sends.



Traceability guardrail

Classification, rationale and confidence are logged before handoff.

HARDEN BEFORE LIVE ROLLOUT



Input + privacy controls

Treat email as untrusted data. Minimize sensitive text and cap input size.



Schema + destination controls

Check required fields and allowed labels. Reject unexpected routing destinations.



Quality + injection checks

Verify unsupported claims. Test hostile instructions, ambiguous topics and missing context.



Operational safeguards


Deduplicate handoffs, quarantine failures and test routing fallback and recovery.

 **Valid JSON can still be wrong. Confidence scores never replace human review.**

Make failures visible, recoverable and testable.

Before live rollout: add a failure path, expand evaluation and keep manual inbox handling available.

FAILURE PATH TO ADD



Retry the failed step, not every completed handoff.
Reconcile sends whose delivery status is uncertain.

TEST EVIDENCE + NEXT GATE

25	tickets reviewed to define request types	15	prototype cases tested
-----------	--	-----------	------------------------

Next: held-out cases for urgency, multi-topic emails, missing history, prompt injection and routing fallback.

Release gate: version model, prompt and scenario; rerun tests.

 Track quality and operations

Wrong owner, urgent misses and human edits.
Errors, oldest unreviewed item, latency and token use.

 Safe replay + manual fallback

Track message ID + destination in a Make Data Store.
Pause the scenario and use the inbox if failures persist.

Test counts: assignment, pp. 1-2. | [Recovery design: Make error handlers](#) | [Retry behavior](#)