

Adoption & Enablement Playbook

How to use, review, train, support, and hand over the email-triage AI workflow

Workflow Gmail → Make.com → LLM classification → JSON → audit log → Slack routing → human-reviewed draft	Primary user Business analyst / onboarding specialist reviewing inbound requests and draft replies
Operating principle Automate repetitive interpretation and routing. Keep customer-facing decisions under human control.	Goal of this playbook Enable a first-time AI user to operate the workflow safely without needing to understand the underlying model.

The reviewer does not need to "trust the AI." The reviewer needs to know what the system is allowed to do, what must be checked, and what to do when the output is uncertain.


What a trained reviewer should be able to do

1	Read the AI decision Understand category, priority, destination, confidence, and draft reply.
2	Check it against the source Use the original email as the source of truth, not the model explanation.
3	Approve, edit, reroute, or escalate Take the correct action without changing the workflow configuration.
4	Capture corrections Record mistakes so the prompt, routing rules, and golden dataset can improve.

QUICK ORIENTATION

1. Understand the workflow before using it

The AI is the interpretation layer. Make.com handles the deterministic workflow around it, while the specialist remains the final control point before any customer-facing reply is sent.



What happens to each message

1 Gmail intake

New email triggers the Make scenario.

2 LLM interpretation

Classify type, priority, route and draft.

3 JSON validation

Convert the response into fixed fields.

4 Audit log

Store the decision before any routing action.

5 Slack routing

Send internally to the responsible team.

6 Human review

Verify the source email before sending.

Human-control boundary

Slack routing is internal. Customer-facing replies stay human-reviewed in the initial rollout.

USE THIS FOR EVERY MESSAGE

2. Reviewer quick start

1

Open the routed Slack item

Read the original email first. Do not start with the AI draft.

2

Check category and owner

Ask: is this actually the right request type, and is this team responsible for the next action?

3

Check priority

Look for deadlines, go-live dates, safety, privacy, service interruption, and business impact. Do not infer urgency from tone alone.

4

Check the draft reply

Verify facts, commitments, dates, names, links, and any statement that sounds like a policy or promise.

5

Choose an action

Approve, edit, reroute, or escalate. If context is missing, request clarification rather than guessing.

6

Record corrections

If the AI was wrong, capture the corrected category, priority, route, or draft issue in the feedback log.

30-second review rule

Source → Decision → Consequence. Confirm what the email actually says, whether the AI decision matches it, and what happens if that decision is wrong.

DECISION GUIDE

3. How to review the AI decision

Check	Approve when...	Edit / reroute when...	Escalate when...
Classification	The primary issue matches the category and no material secondary issue is missed.	The category is wrong or a second topic requires another team.	Ownership is unclear or the request spans multiple teams with conflicting actions.
Priority	Priority reflects deadline, impact, go-live timing, safety, privacy, or service interruption.	Tone has distorted urgency, or a real deadline was missed.	P0/P1, privacy, safety, or legal concerns need immediate human attention.
Routing	The destination team owns the next action.	A different team should act first, or multiple teams need visibility.	No owner exists or the request sits outside the approved taxonomy.
Draft reply	The response is factual, useful, appropriately toned, and does not over-promise.	Facts, wording, next steps, or tone need correction.	The draft contains unsupported policy, sensitive advice, or a commitment the reviewer cannot authorize.

Example A · Calm but urgent	"No rush, but school starts the day after tomorrow and we still cannot add teachers."	Reviewer decision: urgency comes from the go-live deadline, not the polite tone.
Example B · Multi-topic	"Invoice is wrong and the admin account is locked."	Reviewer decision: billing plus product/support. Do not force one category if both require action.

EXCEPTIONS AND FAILURE HANDLING

4. When the workflow should not be trusted

Multi-topic email	More than one actionable issue is present.	Route to every required team or send to specialist review. Confirm one issue is not hiding a more urgent second issue.
Missing conversation history	The email says “following up” or depends on a prior thread the model cannot see.	Do not allow the draft to invent context. Open the thread / CRM record and review manually.
Low confidence or unclear category	The output falls outside the approved taxonomy or is ambiguous.	Treat confidence as a signal, not proof. Send to general review and assign manually.
Privacy, safety, or sensitive content	The email contains personal, regulated, legal, safety, or security-related information.	Use the escalation path. Do not copy content into unapproved tools or channels.
Wrong or unsupported draft	The draft invents facts, policy, dates, product behavior, or commitments.	Delete or rewrite the unsupported content. Log the failure type for prompt / knowledge-base improvement.
Workflow / integration error	A Make step fails, data is malformed, duplicate messages appear, or Slack routing does not complete.	Use the original Gmail inbox as the recovery source. Do not assume a failed automation means the request is being handled elsewhere.

Fail safely

If the AI output is uncertain, incomplete, or difficult to explain, move the request back to the normal human process. The fallback is a feature of the system, not a failure of adoption.

OPERATING HABITS

5. Best practices for first-time AI workflow users

DO

- **Use the source email as ground truth.**
- **Verify high-consequence fields.** Deadlines, privacy, customer commitments, escalation level.
- **Correct the AI when it is wrong.** Corrections improve the system only when they are captured.
- **Keep requests within approved tools.** Gmail, Make, approved model connection, Sheets, and Slack.
- **Use the fallback path.** Ambiguity should increase human review, not force a model decision.

DO NOT

- **Do not send a draft because it “sounds right.”**
- **Do not treat confidence as accuracy.** The model can be confidently wrong.
- **Do not work around guardrails.** If a case is blocked, escalate it.
- **Do not paste customer data into personal AI tools.**
- **Do not silently fix recurring errors.** Log patterns so the workflow can be improved.

A simple mental model

	Classification, priority, destination, draft reply.
	Parsing, logging, routing structure, approved channels, and fallback behavior.
	Whether the decision is appropriate and whether a customer-facing response can be sent.

ENABLEMENT PLAN

6. Training approach used for rollout

Training was structured around the reviewer's actual decisions rather than generic AI concepts. The goal was to build confidence through guided use, supervised practice, and visible support.

Method	Format	What happens
Kickoff session	45-60 min	Walk through the business problem, workflow, what AI can and cannot do, human-review rule, and two live examples.
Role-based 1:1	30 min per user	Use the reviewer's own request types. Practice approve, edit, reroute, and escalate decisions. Capture questions that should become playbook updates.
Supervised practice	First 10-20 cases	User operates the workflow while I observe. Focus on high-risk checks and consistent correction logging, not speed.
Office hours	2x weekly during pilot	Open channel for edge cases, questions, and examples that do not fit the taxonomy. Review recurring confusion and update guidance.
Handover session	60 min	Transfer operating ownership, access map, prompt / routing documentation, feedback log, known limitations, and escalation contacts.
Follow-up review	After 1-2 weeks	Review correction patterns, adoption friction, workload impact, and whether any rule, prompt, or training material needs adjustment.

Training success is behavioral

A user is ready when they can explain the workflow boundary, detect a bad decision, use the fallback path, and capture a correction without assistance. "Understands AI" is not the acceptance criterion.

REVIEWER FEEDBACK LOOP

7. Feedback becomes system improvement

1	2	3	4	5
Reviewer corrects Wrong category, priority, route, or draft issue.	Correction logged Store email ID, AI output, human correction, reason.	Weekly review Group repeated failures by taxonomy.	Targeted change Update prompt, rules, examples, or routing map.	Regression test Run golden-set and failure cases before release.

Minimum feedback record

Field	Example	Why it matters
Email / case ID	MSG-1048	Trace the correction back to the original input.
AI decision	Support · P2 · #support-product	Shows what the system actually produced.
Human correction	Billing · P1 · #sales-billing	Creates the ground-truth correction.
Failure reason	Missed launch deadline in body text	Lets repeated failures be grouped into actionable patterns.
Action taken	Prompt example added + regression case	Connects feedback to a controlled system change.

WHO OWNS WHAT

8. Ownership after handover

Area	Primary owner	Responsibility
Day-to-day review	Business analyst / specialist	Review outputs, send customer replies, correct classifications, use fallback path.
Business taxonomy	Process owner + team leads	Approve categories, priority definitions, ownership changes, and escalation rules.
Workflow configuration	Automation owner	Maintain Make scenario, integrations, routing rules, error handling, and credentials.
Prompt and evaluation	AI product / workflow owner	Version prompts, maintain golden set, review failure patterns, run regression tests.
Data and access	IT / security owner	Approve API access, least-privilege permissions, retention, and approved data handling.
Adoption and documentation	Team lead / enablement owner	Maintain playbook, onboard new reviewers, run refresher training and office hours as needed.

Handover package

Operating This playbook Escalation matrix Channel map	Technical Make scenario ownership Integration / credential map Known failure modes	AI quality Prompt version Golden dataset Evaluation results Regression checklist	Improvement Feedback log Change log Office-hours notes Open backlog
---	--	---	--

Handover is complete when the team can operate without the builder

A named owner can explain the workflow, access the logs, recover from a failed run, train a new reviewer, and make a controlled change through the evaluation and regression process.

KEEP BESIDE THE WORKFLOW

9. One-page reviewer desk card

BEFORE YOU APPROVE	IF SOMETHING LOOKS WRONG
• Read the original email. • Check category and owner. • Check urgency from facts, not tone. • Verify draft facts and commitments. • Confirm no secondary issue was missed.	• Do not force a decision. • Use specialist / general review. • Open prior thread if context is missing. • Escalate privacy, safety, P0/P1. • Log the correction and reason.

Four allowed reviewer actions

APPROVE Decision is correct and draft is safe to send after normal review.	EDIT Core route is correct, but wording or details need correction.
REROUTE Wrong owner or multiple teams need the request.	ESCALATE High-risk, ambiguous, unsupported, or outside the approved taxonomy.

Final rule

If you cannot explain why the AI decision is safe, do not send it. Return the case to human handling, capture what confused the system, and use that evidence to improve the workflow.

Source note

Workflow details are based on the documented prototype: Gmail intake, Make.com orchestration, LLM classification and draft generation, JSON parsing, Google Sheets audit logging, Slack routing, and human review before any customer response. The prototype also identified multi-topic requests, missing conversation history, and limited evaluation data as key failure modes.