


From prototype to improvement decisions

I used the email-triage workflow to show what changed, what the team adopted and where the next investment should go.

Built on prior steps: discovery selected support-email triage, evaluation tested it, and operations made it monitorable.



Positioning note: this is not a fake case study. It is the operating model I used to prove value and decide the next workflow responsibly.

Measured results: simple operating scorecard

I separated prototype evidence from live-pilot metrics so the outcome story stayed honest and useful.

Prototype evidence

Reviewed source tickets

25

Taxonomy + request types

Tested workflow cases

15

Against expected labels

Audit mechanism

Sheets + Slack

Decision log + reviewer queue

I did not treat the prototype as production proof. I used it to create the measurement structure for a controlled pilot.

Live pilot metrics I would review

Outcome = capacity + trust + reliability + cost. Accuracy alone was not enough.

Metric	What I checked	Decision signal
Handling effort	Active minutes per email, incl. review	Is time actually recovered?
Routing quality	Correct owner/channel and override reason	Can teams trust the handoff?
Urgent recall	P0/P1 cases caught; no hidden urgent misses	Is the workflow safe enough?
Review effort	Draft edits, rejections and low-confidence cases	Where does prompt/rule tuning help?
Cost control	Make credits + OpenAI token spend per run/day	Is it affordable to operate?

Result I wanted to show after pilot: fewer minutes spent sorting email, fewer missed urgent cases, lower rework and a visible audit trail for every AI suggestion.

The numbers on this page are outcome categories and measurement structure; final production results would be filled from the live pilot dashboard.

Adoption, benefits and lessons learned

The real benefit was not just automation. It was making intake quicker, reviewable and easier to improve.

Business benefits

- Reduced repetitive reading and sorting
- Faster route to the right owner
- Reusable intake pattern for other queues

Adoption signals

- Specialist used the review queue
- Teams received clearer context
- Corrections became improvement inputs

Lessons learned

- Rules must come before AI judgment
- Low-confidence cases need own path
- Small prompt changes need retesting

Improvement loop



Key learning: an AI workflow becomes useful when users can correct it, see why it acted, and trust that bad changes can be rolled back.

Next investment priority: Scrum Master ticket drafting

After email triage, I selected the highest-ranked next workflow and kept the rollout narrow and reversible.

Ranked investment choice

Workflow	Score	Decision
Business analyst email triage	3.80	Operate + improve
Scrum Master ticket drafts	3.45	Next build
Document intake	2.90	Prepare evidence
Architecture request routing	2.55	Mature owner map

Selected next: Scrum Master ticket drafting

Why Scrum Master next

- Same pattern as email triage: unstructured input to reviewable output
- Human approval keeps risk low before Jira creation
- Good reuse: prompt library, golden-set testing, correction logs

Possible workflows

- Meeting notes -> Jira drafts
- Slack decisions -> ticket candidates
- Backlog refinement -> missing-field list

Why not the others first


- Architecture routing has high value, but owner boundaries were not mature enough.
- Document intake needs more examples across invoices, formats, duplicates and missing references.
- Ticket drafting had the clearest next path from validated email-triage practices.

Investment rule: do not scale AI into unclear ownership. Fix the workflow rules first, then automate the repeatable part.


The ranking uses the discovery scorecard: email triage was first, ticket drafting was the strongest next candidate after that.

Roadmap for the next workflow: ticket drafting

Goal: reduce manual transcription while preventing invented owners, dates or acceptance criteria.



Future-state flow



Success measures

- Time to accepted ticket
- Required-field completeness
- Correction and duplicate rate

Control boundary

- No invented owner, dates or commitments
- Open fields stay explicit
- Human approval before Jira


Reuse from email triage

- Prompt versioning
- Golden-set checks
- Review correction log

Investment roadmap: one pattern, reused carefully

The next stage was not “more AI everywhere.” It was controlled expansion into workflows with the same repeatable pattern.

Roadmap logic: expand from proven pattern



Investment principle: scale only where five gates are true - bounded workflow, available examples, named reviewer, audit log and manual fallback.

Portfolio message: I can move from discovery to implementation, then keep improving workflows with evidence instead of assumptions.